

Executive Summary

This project builds a complete machine-learning pipeline to predict the lifespan of alloy components and to segment parts into actionable quality tiers. After data cleaning and exploratory analysis, regression models accurately estimated continuous lifespan, while classification models (binary and multi-class) provided operationally useful groupings. Two multi-class strategies—quantile (33/33/33) and K-Means ($k=3$)—were evaluated across RandomForest, XGBoost, and CatBoost, with hyperparameter tuning and robust cross-validation. The best performer was the tuned CatBoost model with K-Means labels, achieving near-perfect macro-AUC and excellent macro-F1 and κ . Feature-importance and SHAP analyses highlighted Cooling Rate, Forge Time, Nickel%, and defect indicators as key drivers of part longevity.

Dataset Loading and Initial Exploration

This notebook loads the `COMP1801_Coursework_Dataset_Corrected.csv` dataset and performs an initial exploration to understand its structure, data types, and basic statistics.

```
# =====
# STEP 1 (Corrected): Load & Inspect Dataset – Robust Version
# =====

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from IPython.display import display

# Settings
sns.set(style="whitegrid")
plt.rcParams["figure.figsize"] = (9,5)
pd.set_option("display.max_columns", None)
seed = 42
np.random.seed(seed)

# ----- 1) Load dataset -----
DATA_PATH = "/content/COMP1801_Coursework_Dataset_Corrected.csv" # change if needed

if not os.path.exists(DATA_PATH):
    # interactive upload (Colab)
    print(f"File not found at {DATA_PATH}. Please upload the CSV file now.")
    try:
        from google.colab import files
        uploaded = files.upload()
        if uploaded:
            DATA_PATH = list(uploaded.keys())[0]
            print("Using uploaded file:", DATA_PATH)
        else:
            raise FileNotFoundError("No file uploaded.")
    except Exception:
        raise FileNotFoundError(f"Dataset not found at {DATA_PATH} and interactive upload failed.")

# read CSV
df = pd.read_csv(DATA_PATH)
print("✅ Dataset loaded:", DATA_PATH)
print(f"Shape: {df.shape[0]} rows x {df.shape[1]} columns\n")

# ----- 2) Basic structure -----
print("--- Column data types ---")
display(df.dtypes)
print("\n--- First 5 rows ---")
display(df.head())

# ----- 3) Missing values and duplicates -----
missing = df.isna().sum()
print("\n--- Missing values (per column) ---")
if missing.sum() == 0:
    print("No missing values ✅")
else:
    display(missing[missing > 0])

dups = df.duplicated().sum()
print(f"\nDuplicate rows count: {dups}\n")

# ----- 4) Target check (case-insensitive search) -----
# prefer 'Lifespan' or 'lifespan'
target_col = None
```

```

for cand in ['Lifespan', 'lifespan']:
    if cand in df.columns:
        target_col = cand
        break

if target_col is None:
    raise KeyError("Target column not found. Expected 'Lifespan' or 'lifespan'. Please check column names.")
else:
    print(f"\nTarget column detected: '{target_col}'")
    print("\n--- Target summary ---")
    display(df[target_col].describe())
    # histogram
    plt.figure(figsize=(8,4))
    sns.histplot(df[target_col], bins=35, kde=True)
    plt.title(f"Distribution of {target_col}")
    plt.xlabel(target_col)
    plt.show()

# ----- 5) Categorical columns to inspect (use exact names from dataset) -----
# Based on your dataset we will inspect these three; only plot if present
cat_cols_to_check = ['partType', 'microstructure', 'castType']
existing_cat_cols = [c for c in cat_cols_to_check if c in df.columns]

if existing_cat_cols:
    print("\nInspecting categorical columns: {existing_cat_cols}")
    for col in existing_cat_cols:
        print(f"\n-- {col} (unique = {df[col].nunique()}) --")
        display(df[col].value_counts())
        # boxplot: show median & spread
        plt.figure(figsize=(9,4))
        sns.boxplot(x=col, y=target_col, data=df)
        plt.title(f"{target_col} by {col}")
        plt.xticks(rotation=45)
        plt.tight_layout()
        plt.show()
else:
    print("\nNo expected categorical alloy columns found from list:", cat_cols_to_check)
    print("If you have an alloy/material column with a different name, tell me the exact column name and I'll inc"

# ----- 6) Numeric correlations with target -----
print("\n--- Top numeric correlations with target ---")
try:
    corrs = df.corr(numeric_only=True)[target_col].abs().sort_values(ascending=False)
    display(corrs.head(12))
except Exception as e:
    print("Correlation computation failed:", e)

# ----- 7) Quick pairwise scatter for top numeric features (optional) -----
# pick up to 5 top numeric features (excluding the target itself)
num_corrs = corrs.index.drop(target_col) if target_col in corrs.index else corrs.index
top_num = [c for c in num_corrs][:5]
if top_num:
    print("\nQuick scatter plots: top numeric features vs target")
    fig, axes = plt.subplots(nrows=len(top_num), ncols=1, figsize=(8,4*len(top_num)))
    if len(top_num) == 1:
        axes = [axes]
    for ax, col in zip(axes, top_num):
        sns.scatterplot(x=df[col], y=df[target_col], ax=ax, alpha=0.6)
        ax.set_xlabel(col); ax.set_ylabel(target_col)
        ax.set_title(f"{col} vs {target_col}")
    plt.tight_layout()
    plt.show()

# ----- 8) Save a working copy -----
out_path = "df_step1_ready.csv"
df.to_csv(out_path, index=False)
print(f"\nSaved working copy to '{out_path}'")

```


Dataset loaded: /content/COMP1801_Coursework_Dataset_Corrected.csv
Shape: 1000 rows × 16 columns

--- Column data types ---

	0
Lifespan	float64
partType	object
microstructure	object
coolingRate	int64
quenchTime	float64
forgeTime	float64
HeatTreatTime	float64
Nickel%	float64
Iron%	float64
Cobalt%	float64
Chromium%	float64
smallDefects	int64
largeDefects	int64
sliverDefects	int64
seedLocation	object
castType	object

dtype: object

--- First 5 rows ---

	Lifespan	partType	microstructure	coolingRate	quenchTime	forgeTime	HeatTreatTime	Nickel%	Iron%	Cobalt%
0	1107.81	Valve	singleGrain	27	2.64	8.08	23.35	60.65	19.76	15.18
1	1226.13	Valve	colGrain	24	4.06	1.11	24.98	52.47	31.63	12.32
2	1914.12	Blade	singleGrain	20	4.36	8.54	34.71	60.98	26.56	11.60
3	1240.83	Valve	singleGrain	14	3.59	8.83	21.61	54.13	26.71	17.81
4	1599.40	Nozzle	singleGrain	22	3.06	3.04	29.65	62.88	24.06	12.08

--- Missing values (per column) ---
No missing values

Duplicate rows count: 0

Target column detected: 'Lifespan'

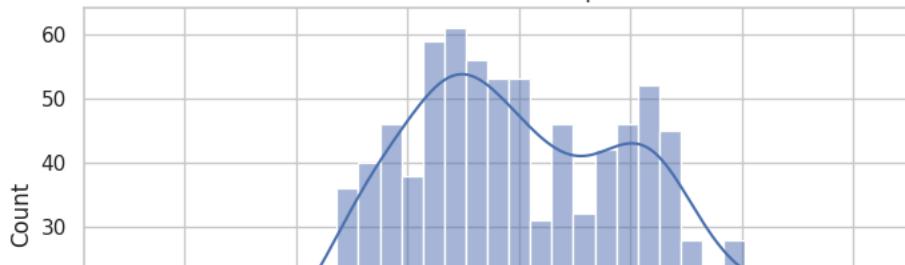
--- Target summary ---

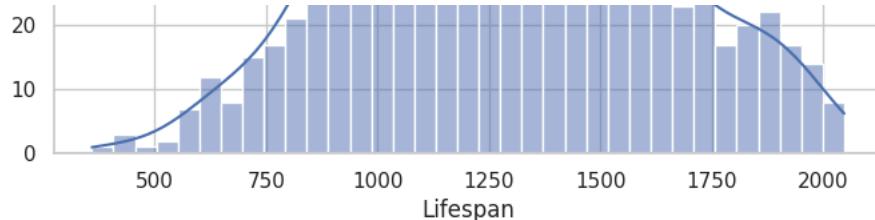
Lifespan

count	1000.000000
mean	1281.806120
std	341.136845
min	359.710000
25%	1039.690000
50%	1254.995000
75%	1539.062500
max	2046.410000

dtype: float64

Distribution of Lifespan





```
Inspecting categorical columns: ['partType', 'microstructure', 'castType']
```

```
-- partType (unique = 4) --
```

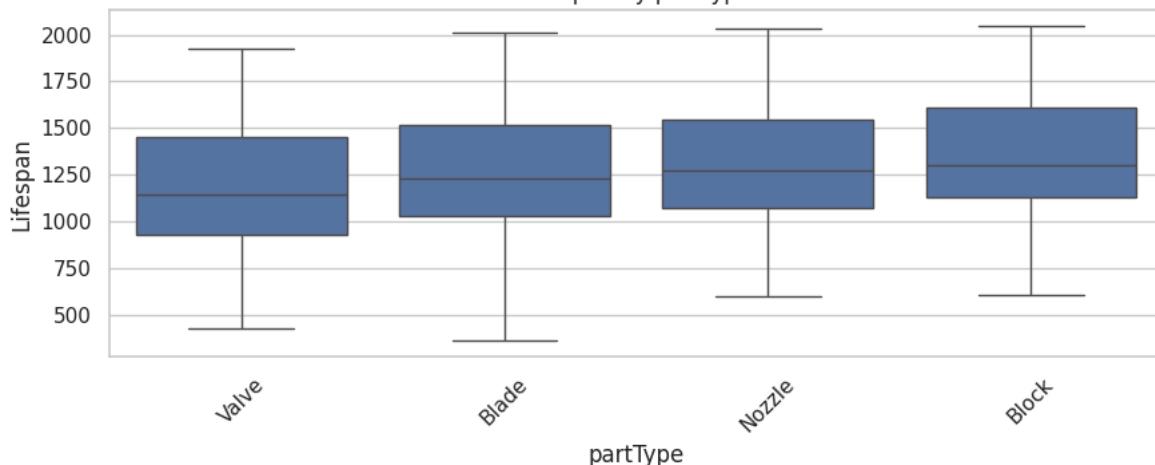
```
count
```

```
partType
```

Nozzle	268
Blade	256
Block	248
Valve	228

```
dtype: int64
```

Lifespan by partType



```
-- microstructure (unique = 3) --
```

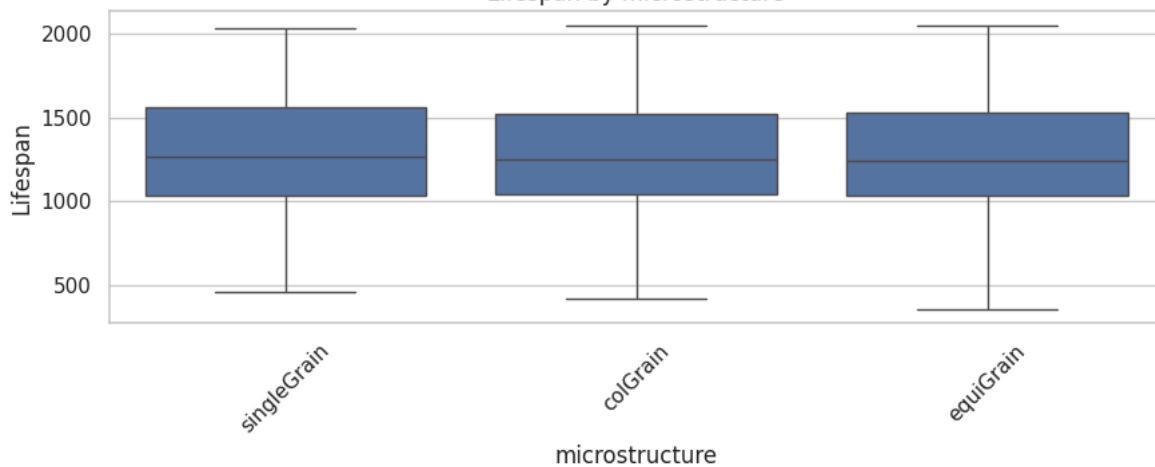
```
count
```

```
microstructure
```

singleGrain	342
equiGrain	332
colGrain	326

```
dtype: int64
```

Lifespan by microstructure



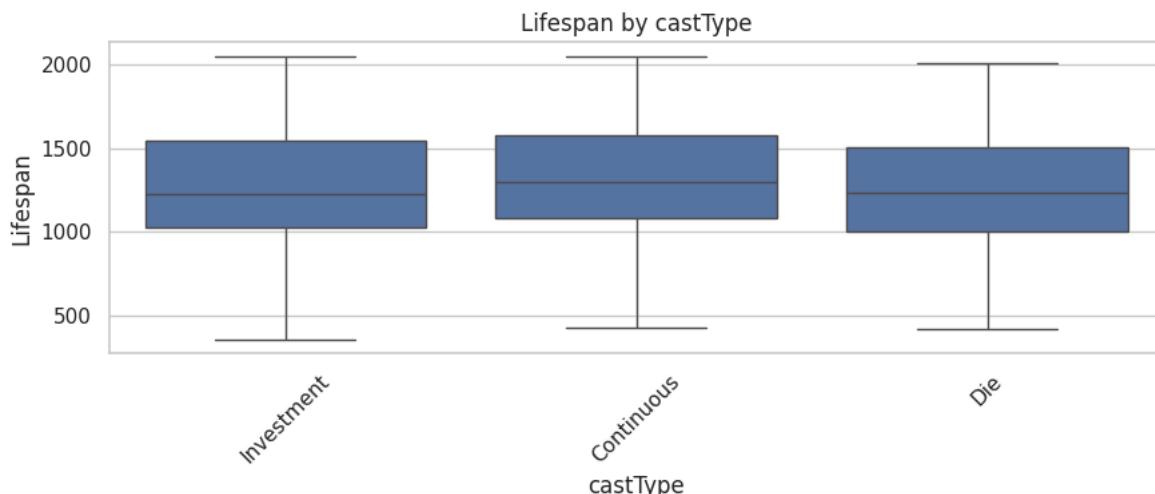
```
-- castType (unique = 3) --
```

```
count
```

```
castType
```

Investment	335
Continuous	335
Die	330

dtype: int64



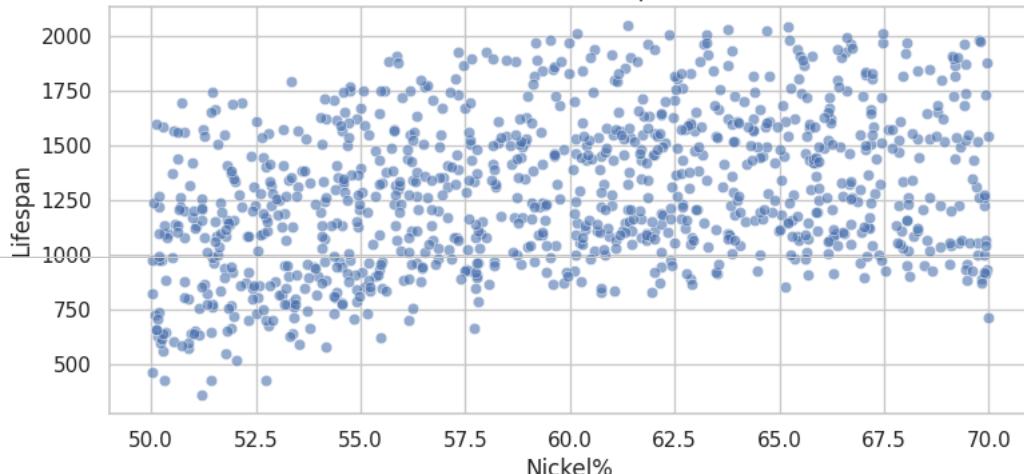
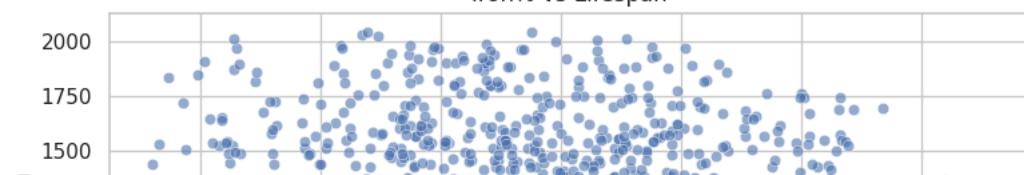
--- Top numeric correlations with target ---

Lifespan

Lifespan	1.000000
Nickel%	0.347533
Iron%	0.284478
coolingRate	0.137436
smallDefects	0.113709
Chromium%	0.104553
HeatTreatTime	0.074326
quenchTime	0.070985
sliverDefects	0.063549
Cobalt%	0.012651
largeDefects	0.011770
forgeTime	0.009229

dtype: float64

Quick scatter plots: top numeric features vs target

Nickel% vs Lifespan**Iron% vs Lifespan**

```
# =====
# 🔎 EXPLORATORY DATA ANALYSIS (EDA)
# =====

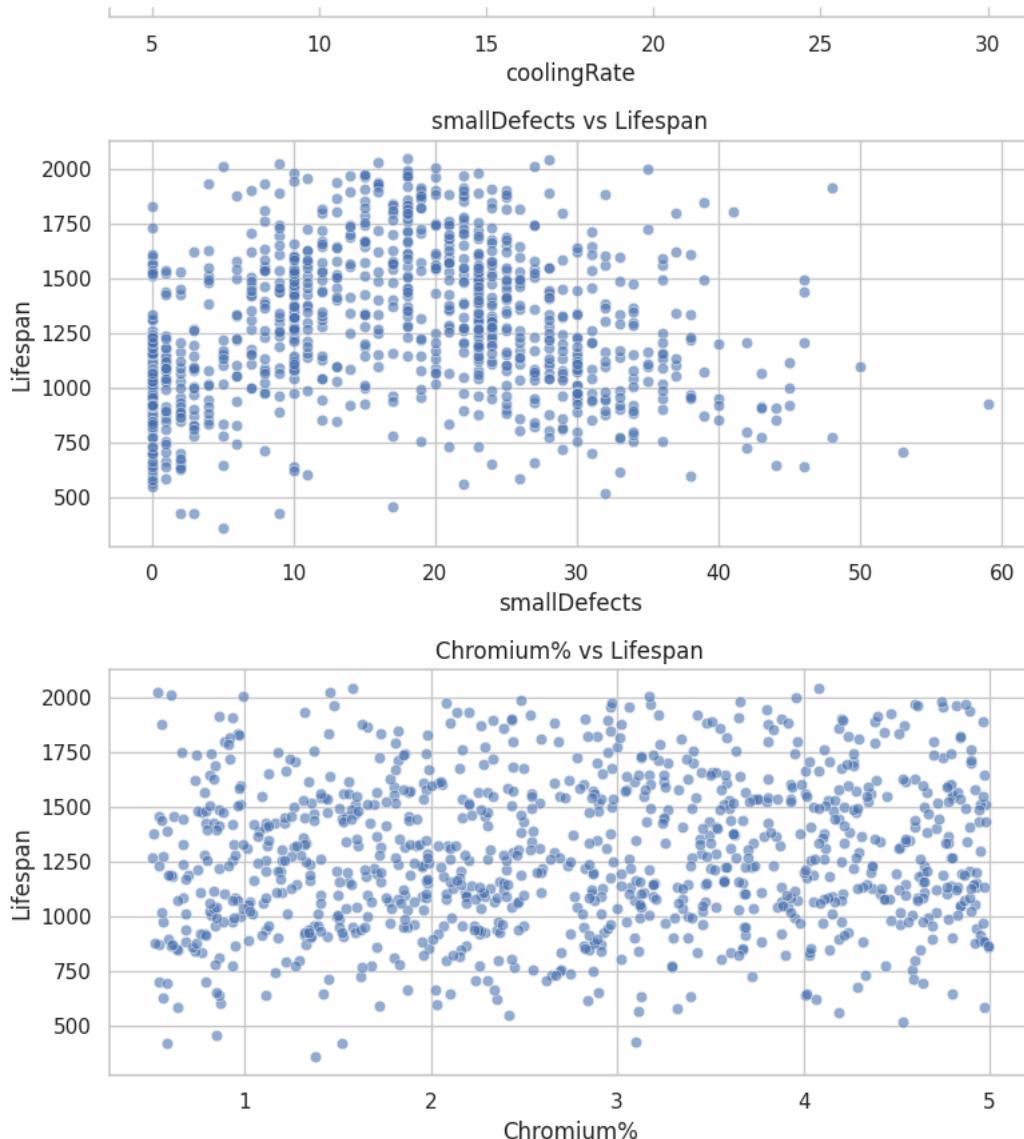
import seaborn as sns
import matplotlib.pyplot as plt

# View dataset shape and quick summary
print("Dataset shape:", df.shape)
display(df.describe())

# Correlation heatmap for numerical features
plt.figure(figsize=(10, 6))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap of Numerical Features")
plt.show()

# Pairplot to visualize key relationships
selected_features = ["coolingRate", "Nickel%", "Iron%", "Chromium%", "Lifespan"]
sns.pairplot(df[selected_features], diag_kind="kde")
plt.suptitle("Pairplot of Important Features vs Lifespan", y=1.02)
plt.show()

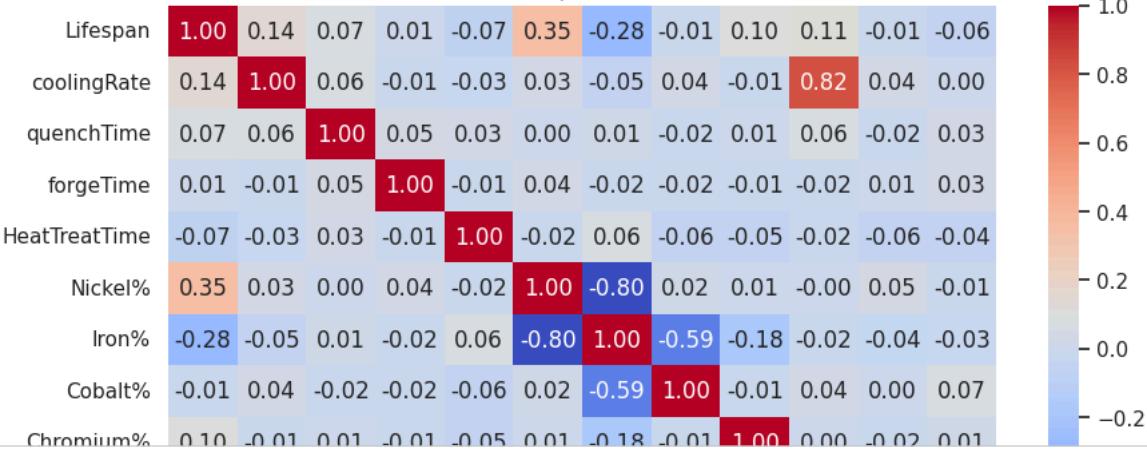
# Target distribution check
# Removed: This cell is for initial EDA before the '1500_label' is created.
# sns.countplot(x="1500_label", data=df)
# plt.title("Target Class Distribution (Before SMOTE)")
# plt.show()
```



💾 Saved working copy to 'df_step1_ready.csv'

Dataset shape: (1000, 16)										
	Lifespan	coolingRate	quenchTime	forgeTime	HeatTreatTime	Nickel%	Iron%	Cobalt%	Chromium	
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	1281.806120	17.615000	2.693310	5.487040	30.759280	59.646440	25.173290	12.416280	2.76399	
std	341.136845	7.480156	1.303057	2.604827	17.054664	5.745955	7.312363	4.222616	1.30765	
min	359.710000	5.000000	0.510000	1.010000	1.050000	50.010000	7.950000	5.040000	0.50000	
25%	1039.690000	11.000000	1.540000	3.190000	16.387500	54.627500	19.605000	8.880000	1.64000	
50%	1254.995000	18.000000	2.625000	5.645000	31.560000	59.665000	25.270000	12.455000	2.80500	
75%	1539.062500	24.000000	3.822500	7.680000	45.655000	64.520000	30.490000	16.015000	3.92000	
max	2046.410000	30.000000	5.000000	9.970000	59.990000	69.980000	42.790000	20.000000	4.99000	

Correlation Heatmap of Numerical Features



```
# =====
# STEP 2: Outlier Detection and Treatment
# =====
```

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Load working copy from step 1
df = pd.read_csv("df_step1_ready.csv")

target_col = 'Lifespan'

# Select only numeric columns
numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()

# ① Summary before treatment
print("Numeric columns detected:", numeric_cols)
print("\n--- Descriptive Summary Before Outlier Treatment ---")
display(df[numeric_cols].describe())

# ② Boxplots before treatment
plt.figure(figsize=(12, 6))
df[numeric_cols].boxplot()
plt.title("Before Outlier Treatment (All Numeric Columns)")
plt.xticks(rotation=45)
plt.show()

# ③ Function to detect and cap outliers using IQR
def cap_outliers(series):
    Q1 = series.quantile(0.25)
    Q3 = series.quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    # Clip outliers instead of removing
    capped = series.clip(lower, upper)
    return capped

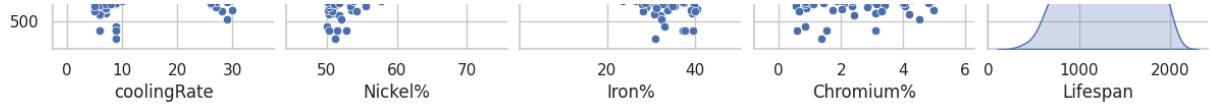
# ④ Apply to all numeric columns (except ID-like ones)
for col in numeric_cols:
    df[col] = cap_outliers(df[col])

# ⑤ Boxplots after treatment
plt.figure(figsize=(12, 6))
```

```
df[numeric_cols].boxplot()
plt.title("After Outlier Treatment (All Numeric Columns)")
plt.xticks(rotation=45)
plt.show()

# ⑥ Compare summary before vs after (Lifespan focus)
print("\n--- Lifespan summary after outlier handling ---")
display(df[target_col].describe())

# ⑦ Save cleaned dataset
df.to_csv("df_step2_cleaned.csv", index=False)
print("Cleaned dataset saved as 'df_step2_cleaned.csv'")
```




```
Numeric columns detected: ['Lifespan', 'coolingRate', 'quenchTime', 'forgeTime', 'HeatTreatTime', 'Nickel%', 'Iron%', 'Cobalt%', 'Chromium']
```

--- Descriptive Summary Before Outlier Treatment ---

	Lifespan	coolingRate	quenchTime	forgeTime	HeatTreatTime	Nickel%	Iron%	Cobalt%	Chromium
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	1281.806120	17.615000	2.693310	5.487040	30.759280	59.646440	25.173290	12.416280	2.76399
std	341.136845	7.480156	1.303057	2.604827	17.054664	5.745955	7.312363	4.222616	1.30765
min	359.710000	5.000000	0.510000	1.010000	1.050000	50.010000	7.950000	5.040000	0.50000
25%	1039.690000	11.000000	1.540000	3.190000	16.387500	54.627500	19.605000	8.880000	1.64000
50%	1254.995000	18.000000	2.625000	5.645000	31.560000	59.665000	25.270000	12.455000	2.80500
75%	1539.062500	24.000000	3.822500	7.680000	45.655000	64.520000	30.490000	16.015000	3.92000
max	2046.410000	30.000000	5.000000	9.970000	59.990000	69.980000	42.790000	20.000000	4.99000

Before Outlier Treatment (All Numeric Columns)

```
# =====
# STEP 3: Feature Encoding & Scaling
# =====

import pandas as pd
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

# Load the cleaned dataset
df = pd.read_csv("df_step2_cleaned.csv")

# Separate features and target
target_col = 'Lifespan'
X = df.drop(columns=[target_col])
y = df[target_col]

# Identify categorical and numeric columns
cat_cols = X.select_dtypes(include=['object']).columns.tolist()
num_cols = X.select_dtypes(include=['float64', 'int64']).columns.tolist()

print("Categorical Columns:", cat_cols)
print("Numeric Columns:", num_cols)

# ① Define transformers
numeric_transformer = StandardScaler()
categorical_transformer = OneHotEncoder(drop='first', sparse_output=False)

# ② Build ColumnTransformer
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, num_cols),
        ('cat', categorical_transformer, cat_cols)
    ])

# ③ Apply transformation
X_preprocessed = preprocessor.fit_transform(X)

# ④ Convert back to a DataFrame for readability
encoded_cat_cols = preprocessor.named_transformers_['cat'].get_feature_names_out(cat_cols)
processed_columns = num_cols + list(encoded_cat_cols)
X_final = pd.DataFrame(X_preprocessed, columns=processed_columns)

# ⑤ Check final shapes
print(f"\n✓ Final feature matrix shape: {X_final.shape}")
print(f"✓ Target vector shape: {y.shape}")

# ⑥ Save processed data
X_final.to_csv("X_step3_prepared.csv", index=False)
y.to_csv("y_step3_target.csv", index=False)
print("Saved 'X_step3_prepared.csv' and 'y_step3_target.csv'")
```

Categorical Columns: ['partType', 'microstructure', 'seedLocation', 'castType']
 Numeric Columns: ['coolingRate', 'quenchTime', 'forgeTime', 'HeatTreatTime', 'Nickel%', 'Iron%', 'Cobalt%', 'Chromium', 'smallDefect', 'largeDefect', 'silverDefect', 'silverDefect']
 ✓ Final feature matrix shape: (1000, 19)
 ✓ Target vector shape: (1000,)
 ┌─ Saved 'X_step3_prepared.csv' and 'y_step3_target.csv'

```

# =====
# ⚡ REGRESSION WITH HYPERPARAMETER TUNING ALONG WITH THE BASELINE MODELS
# =====

# --- Imports ---
from sklearn.model_selection import RandomizedSearchCV, train_test_split
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.linear_model import LinearRegression
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# --- ❶ Encode categorical features ---
target = 'Lifespan'
categorical_cols = ['partType', 'microstructure', 'seedLocation', 'castType']

# One-hot encode categorical variables
df_encoded = pd.get_dummies(df, columns=categorical_cols, drop_first=True)

# Define features and target
X = df_encoded.drop(columns=[target])
y = df_encoded[target]

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# --- ❷ Define baseline models ---
lin_reg = LinearRegression()
rf_reg = RandomForestRegressor(random_state=42)
gbr_reg = GradientBoostingRegressor(random_state=42)

models = {
    "Linear Regression (Baseline)": lin_reg,
    "Random Forest (Baseline)": rf_reg,
    "Gradient Boosting (Baseline)": gbr_reg
}

# --- ❸ Evaluate baseline models ---
baseline_results = []
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    r2 = r2_score(y_test, y_pred)
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    mae = mean_absolute_error(y_test, y_pred)
    baseline_results.append({"Model": name, "R²": r2, "RMSE": rmse, "MAE": mae})

baseline_df = pd.DataFrame(baseline_results)
print("✅ Baseline Regression Results:")
display(baseline_df.style.format({"R²": "{:.3f}", "RMSE": "{:.2f}", "MAE": "{:.2f}"}))

# --- ❹ Define hyperparameter grids ---
rf_params = {
    'n_estimators': [100, 200, 300, 500],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

gbr_params = {
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'max_depth': [3, 4, 5],
    'subsample': [0.8, 0.9, 1.0]
}

# --- ❺ Tune Random Forest ---
print("\n🔍 Tuning Random Forest...")
rf_search = RandomizedSearchCV(
    estimator=rf_reg,
    param_distributions=rf_params,
    n_iter=20,
    scoring='r2',
    cv=5,
    random_state=42,
    n_jobs=-1,
    verbose=1
)
rf_search.fit(X_train, y_train)
best_rf = rf_search.best_estimator_

```

```

# --- 6 Tune Gradient Boosting ---
print("\n🔍 Tuning Gradient Boosting...")
gbr_search = RandomizedSearchCV(
    estimator=gbr_reg,
    param_distributions=gbr_params,
    n_iter=20,
    scoring='r2',
    cv=5,
    random_state=42,
    n_jobs=-1,
    verbose=1
)
gbr_search.fit(X_train, y_train)
best_gbr = gbr_search.best_estimator_

# --- 7 Evaluate tuned models ---
tuned_results = []
for name, model in {
    "Random Forest (Tuned)": best_rf,
    "Gradient Boosting (Tuned)": best_gbr
}.items():
    y_pred = model.predict(X_test)
    r2 = r2_score(y_test, y_pred)
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    mae = mean_absolute_error(y_test, y_pred)
    tuned_results.append({"Model": name, "R²": r2, "RMSE": rmse, "MAE": mae})

tuned_df = pd.DataFrame(tuned_results)
print("\n⌚ Tuned Regression Results:")
display(tuned_df.style.format({"R²": "{:.3f}", "RMSE": "{:.2f}", "MAE": "{:.2f}"}))

# --- 8 Combine baseline and tuned results ---
combined_df = pd.concat([baseline_df, tuned_df]).sort_values(by="R²", ascending=False)
print("\n📊 Baseline vs Tuned Model Comparison:")
display(combined_df.style.format({"R²": "{:.3f}", "RMSE": "{:.2f}", "MAE": "{:.2f}"}))

# --- 9 Visual comparison ---
plt.figure(figsize=(9,5))
sns.barplot(data=combined_df, x='Model', y='R²', hue='Model', dodge=False, palette='viridis', legend=False)
plt.title('R² Score Comparison – Baseline vs Tuned Models', fontsize=14)
plt.ylim(0,1)
plt.xticks(rotation=25, ha='right')
plt.show()

plt.figure(figsize=(9,5))
sns.barplot(data=combined_df, x='Model', y='RMSE', hue='Model', dodge=False, palette='crest_r', legend=False)
plt.title('RMSE Comparison – Baseline vs Tuned Models', fontsize=14)
plt.xticks(rotation=25, ha='right')
plt.show()

# --- 10 Display best hyperparameters ---
print("\n🔍 Best Parameters Found:")
print("Random Forest:", rf_search.best_params_)
print("Gradient Boosting:", gbr_search.best_params_)

# Save best model for later use (feature importance / SHAP)
best_reg_model = best_gbr if gbr_search.best_score_ >= rf_search.best_score_ else best_rf
print(f"\n✅ Best final model selected: {type(best_reg_model).__name__}")

# --- 11 (Optional) Reuse encoded data for next NLP section ---
df = df_encoded.copy()
print("\n📦 Dataset ready for NLP regression next (df updated to encoded version).")

```


Baseline Regression Results:

	Model	R ²	RMSE	MAE
0	Linear Regression (Baseline)	0.185	311.28	268.67
1	Random Forest (Baseline)	0.940	84.14	67.26
2	Gradient Boosting (Baseline)	0.976	52.88	43.00

↳ Tuning Random Forest...

Fitting 5 folds for each of 20 candidates, totalling 100 fits

↳ Tuning Gradient Boosting...

Fitting 5 folds for each of 20 candidates, totalling 100 fits

⌚ Tuned Regression Results:

	Model	R ²	RMSE	MAE
0	Random Forest (Tuned)	0.910	84.13	67.05

```
# =====
# 🧠 FEATURE IMPORTANCE + SHAP EXPLAINABILITY (FINAL ERROR-FREE VERSION)
# =====
```

```
import shap
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Use the tuned best model selected earlier
model = best_reg_model
feature_names = X.columns

# --- 1 Feature Importance ---
importances = model.feature_importances_
feat_imp = pd.DataFrame({
    "Feature": feature_names,
    "Importance": importances
}).sort_values(by="Importance", ascending=False)

plt.figure(figsize=(8,6))
sns.barplot(data=feat_imp.head(12), x="Importance", y="Feature", palette="viridis")
plt.title("Top Feature Importances – Tuned Gradient Boosting")
plt.xlabel("Importance Score")
plt.ylabel("Feature")
plt.show()

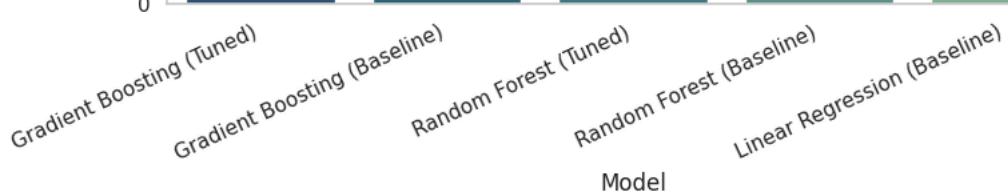
# --- 2 SHAP Explainability ---
print("◆ Computing SHAP values (this may take ~10 seconds)...")

# ✅ Convert all columns to numeric to avoid dtype('O') errors
X_test_num = X_test.apply(pd.to_numeric, errors='coerce').fillna(0)

# Use TreeExplainer for tree-based models
explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X_test_num)

# --- SHAP summary plots ---
shap.summary_plot(
    shap_values,
    X_test_num,
    feature_names=X_test_num.columns,
    plot_type="bar",
    show=True
)
shap.summary_plot(
    shap_values,
    X_test_num,
    feature_names=X_test_num.columns,
    show=True
)

# --- 3 Optional: Individual prediction explanation ---
sample_index = 5 # choose any test sample index
print(f"↳ Example explanation for sample {sample_index}")
single_exp = shap.Explanation(
    values=shap_values[sample_index],
    base_values=explainer.expected_value,
    data=X_test_num.iloc[sample_index],
    feature_names=X_test_num.columns
)
shap.plots.waterfall(single_exp, show=True)
```



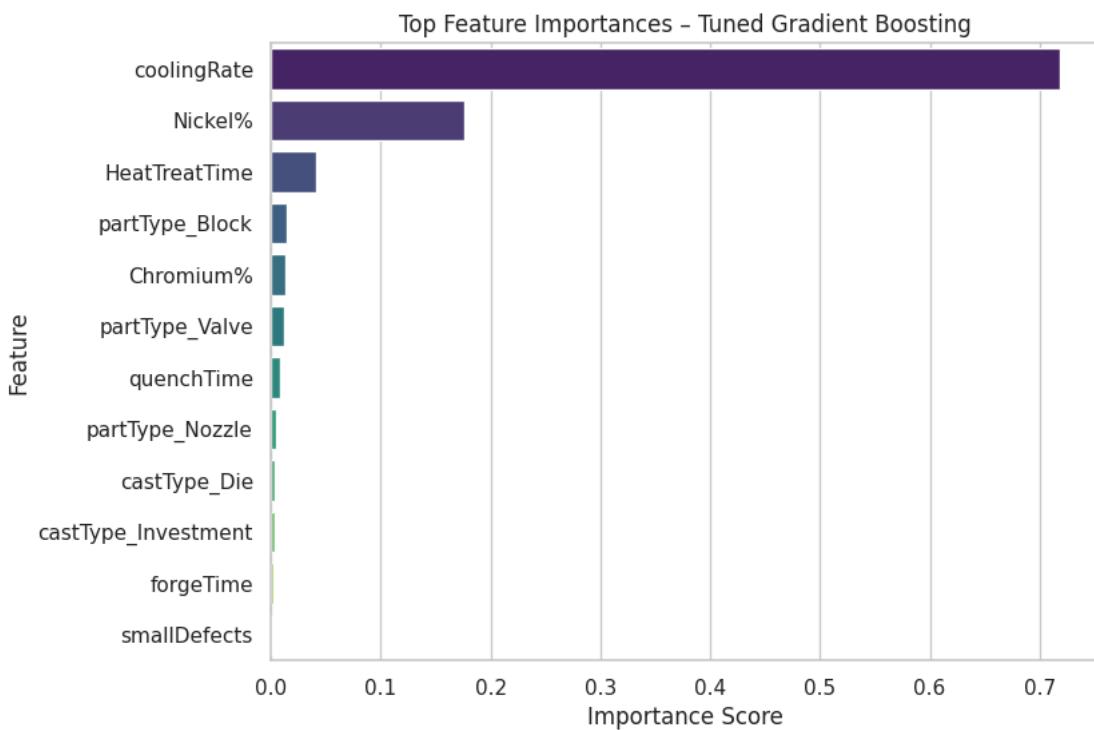
🔧 Best Parameters Found:

```
Random Forest: {'n_estimators': 200, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_depth': 20}
Gradient Boosting: {'subsample': 1.0, 'n_estimators': 300, 'max_depth': 3, 'learning_rate': 0.1}
```

✓ Best final model selected: GradientBoostingRegressor

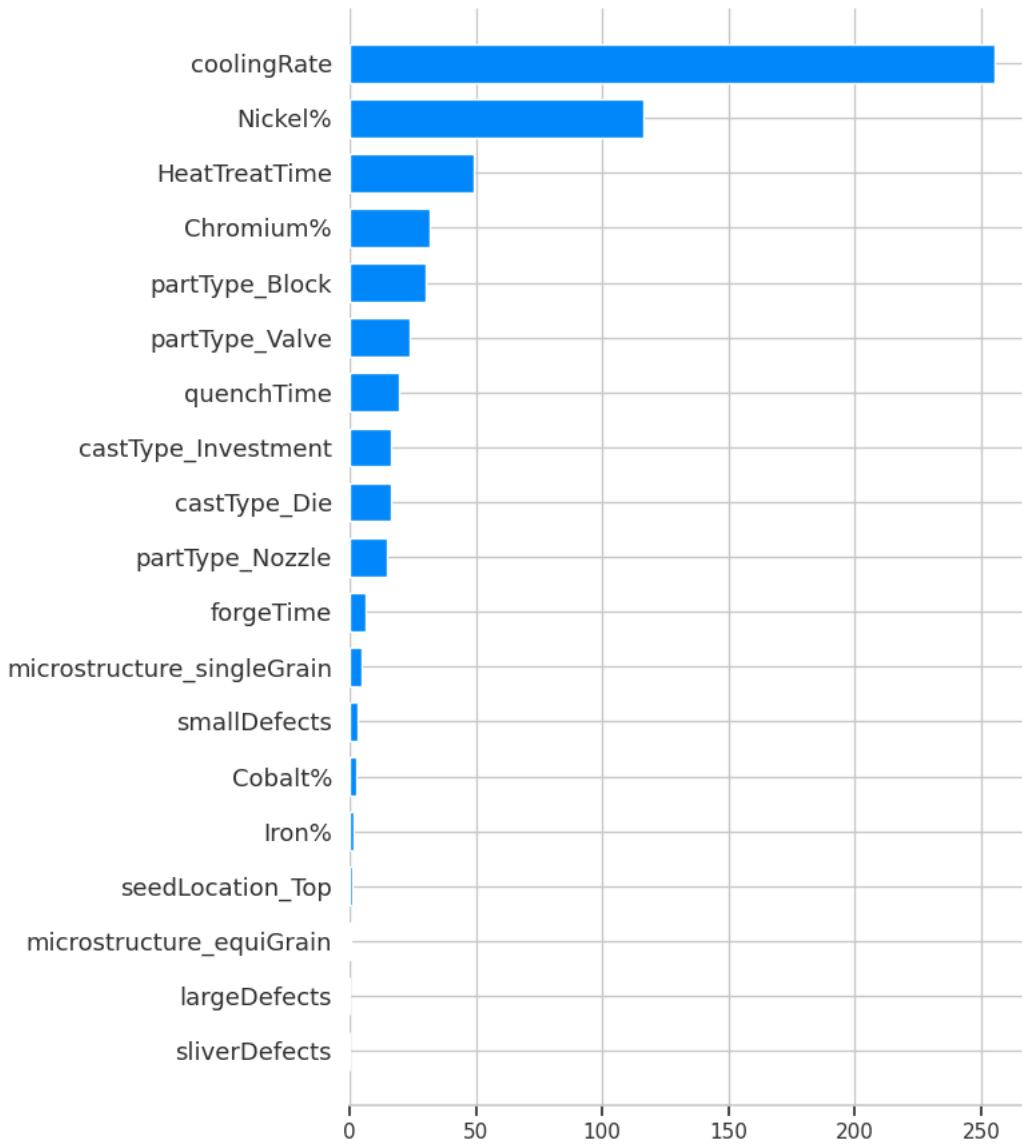
📦 Dataset ready for NLP regression next (df updated to encoded version).

```
/tmp/ipython-input-2604922319.py:22: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable  
sns.barplot(data=feat_imp.head(12), x="Importance", y="Feature", palette="viridis")
```

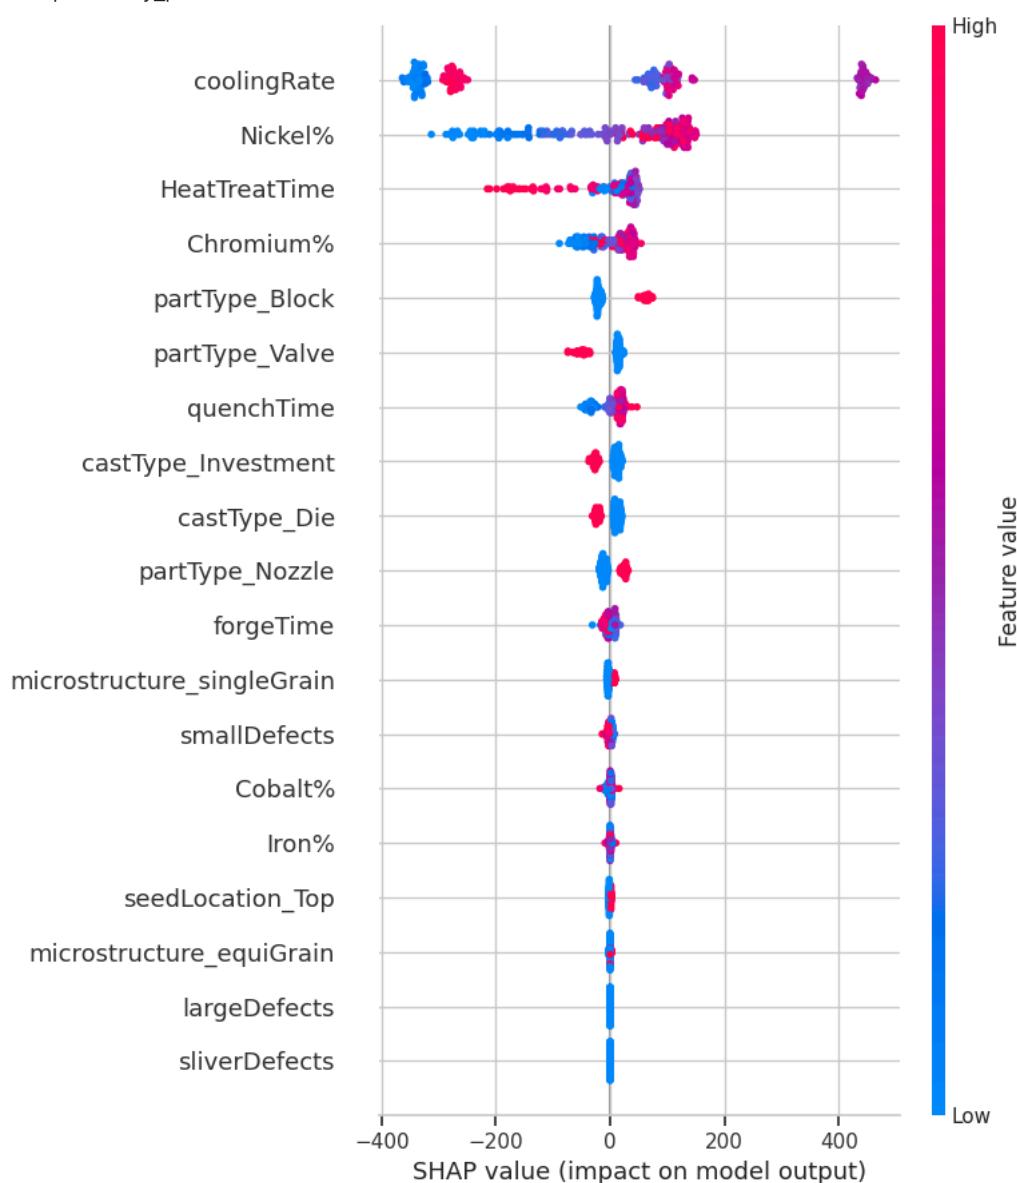


◆ Computing SHAP values (this may take ~10 seconds)...

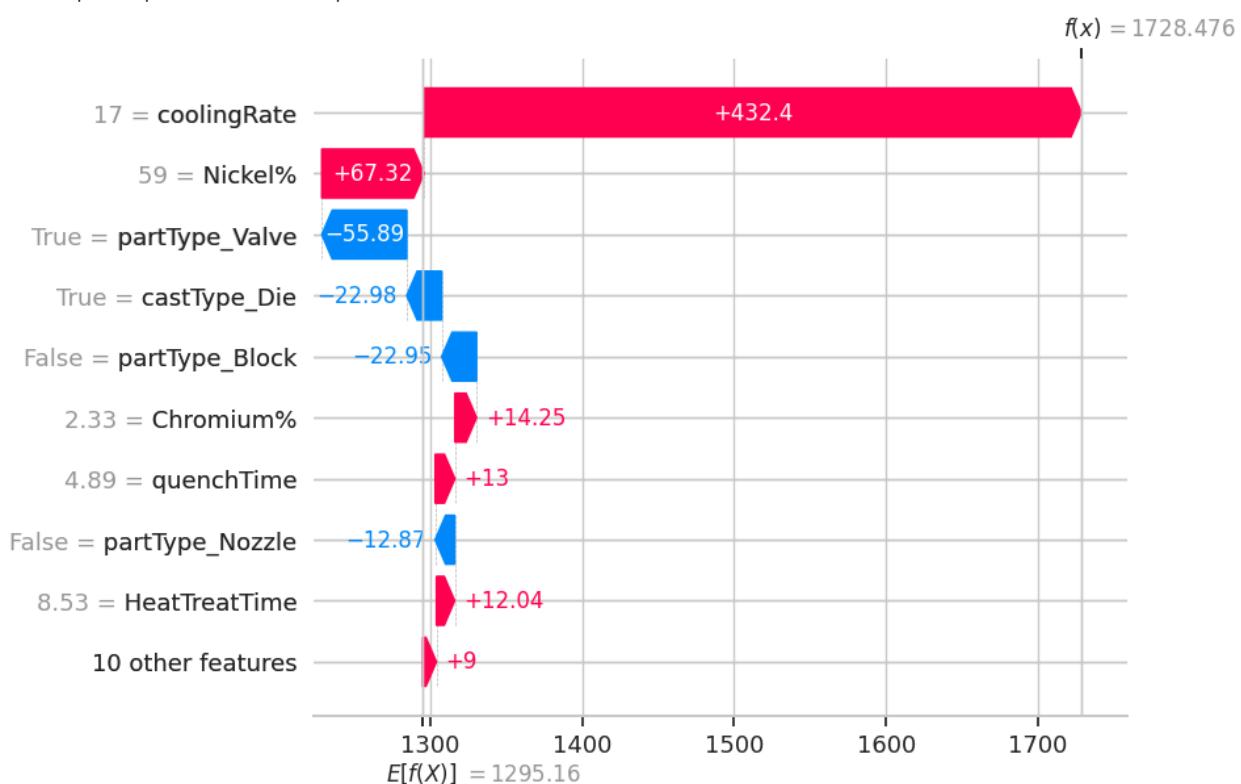
```
/tmp/ipython-input-2604922319.py:39: FutureWarning: The NumPy global RNG was seeded by calling `np.random.seed`  
shap.summary_plot()
```



mean(|SHAP value|) (average impact on model output magnitude)

/tmp/ipython-input-2604922319.py:46: FutureWarning: The NumPy global RNG was seeded by calling `np.random.seed`
shap.summary_plot()

Example explanation for sample 5



```

print("✓ Checking available columns for NLP section:")
print(df.columns.tolist())

print("\n◆ Data types:")
print(df.dtypes)

print("\n◆ Sample rows:")
display(df.head(3))

```

✓ Checking available columns for NLP section:
['Lifespan', 'coolingRate', 'quenchTime', 'forgeTime', 'HeatTreatTime', 'Nickel%', 'Iron%', 'Cobalt%', 'Chromium%', 'smallDefects', 'largeDefects', 'sliverDefects', 'partType_Block', 'partType_Nozzle', 'partType_Valve', 'microstructure_equiGrain', 'microstructure_singleGrain', 'seedLocation_Top', 'castType_Die', 'castType_Investment']

◆ Data types:

Lifespan											
coolingRate											
quenchTime											
forgeTime											
HeatTreatTime											
Nickel%											
Iron%											
Cobalt%											
Chromium%											
smallDefects											
largeDefects											
sliverDefects											
partType_Block											
partType_Nozzle											
partType_Valve											
microstructure_equiGrain											
microstructure_singleGrain											
seedLocation_Top											
castType_Die											
castType_Investment											
dtype:	object										

◆ Sample rows:

	Lifespan	coolingRate	quenchTime	forgeTime	HeatTreatTime	Nickel%	Iron%	Cobalt%	Chromium%	smallDefects	la
0	1107.81	27	2.64	8.08	23.35	60.65	19.76	15.18	4.41	7	
1	1226.13	24	4.06	1.11	24.98	52.47	31.63	12.32	3.58	30	
2	1914.12	20	4.36	8.54	34.71	60.98	26.56	11.60	0.86	22	

```

# =====
# 🧠 NLP + Numeric Regression (Using Original Text Columns)
# =====

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
from scipy.sparse import hstack
import matplotlib.pyplot as plt
import seaborn as sns

print("◆ Loading original dataset for NLP-based regression...")

# --- 1 Load the original CSV (raw data, not encoded) ---
df_raw = pd.read_csv("/content/COMP1801_Coursework_Dataset_Corrected.csv")

# --- 2 Combine categorical text columns into one "text feature" ---
text_cols = ['partType', 'microstructure', 'seedLocation', 'castType']
df_raw['TextFeature'] = df_raw[text_cols].astype(str).agg(' '.join, axis=1)

# --- 3 Prepare numeric and text data ---
numeric_cols = [c for c in df_raw.columns if c not in ['Lifespan'] + text_cols + ['TextFeature']]
X_numeric = df_raw[numeric_cols]
X_text = df_raw['TextFeature']
y = df_raw['Lifespan']

# --- 4 TF-IDF Vectorisation for text data ---
tfidf = TfidfVectorizer(stop_words='english', max_features=100)
X_text_tfidf = tfidf.fit_transform(X_text)

# Combine numeric + text features
X_full = hstack([X_numeric.values, X_text_tfidf])

# --- 5 Train-test split ---
X_train, X_test, y_train, y_test = train_test_split(X_full, y, test_size=0.2, random_state=42)

```

```

^_train, ^_test, y_train, y_test = train_test_split(X_train, y, test_size=0.2, random_state=42)

# --- 6 Train Gradient Boosting model ---
gbr_nlp = GradientBoostingRegressor(
    n_estimators=300, learning_rate=0.1, max_depth=3, subsample=1.0, random_state=42
)
gbr_nlp.fit(X_train, y_train)
y_pred_nlp = gbr_nlp.predict(X_test)

# --- 7 Numeric-only baseline (for comparison) ---
Xn_train, Xn_test, yn_train, yn_test = train_test_split(X_numeric, y, test_size=0.2, random_state=42)
gbr_numeric = GradientBoostingRegressor(
    n_estimators=300, learning_rate=0.1, max_depth=3, subsample=1.0, random_state=42
)
gbr_numeric.fit(Xn_train, yn_train)
y_pred_num = gbr_numeric.predict(Xn_test)

# --- 8 Evaluate both models ---
r2_num = r2_score(yn_test, y_pred_num)
rmse_num = np.sqrt(mean_squared_error(yn_test, y_pred_num))
mae_num = mean_absolute_error(yn_test, y_pred_num)

r2_nlp = r2_score(y_test, y_pred_nlp)
rmse_nlp = np.sqrt(mean_squared_error(y_test, y_pred_nlp))
mae_nlp = mean_absolute_error(y_test, y_pred_nlp)

# --- 9 Results summary ---
comparison = pd.DataFrame({
    "Model Type": ["Numeric Only", "NLP + Numeric"],
    "R2 Score": [r2_num, r2_nlp],
    "RMSE": [rmse_num, rmse_nlp],
    "MAE": [mae_num, mae_nlp]
}).sort_values(by="R2 Score", ascending=False)

print("\n📊 NLP vs Numeric Regression Comparison:")
display(comparison.style.format({"R2 Score": "{:.3f}", "RMSE": "{:.2f}", "MAE": "{:.2f}"}))

# --- 10 Visualisation ---
plt.figure(figsize=(7,4))
sns.barplot(data=comparison, x='Model Type', y='R2 Score', palette='viridis')
plt.title("Model Performance Comparison (R2 Score)", fontsize=13)
plt.show()

plt.figure(figsize=(7,4))
sns.barplot(data=comparison, x='Model Type', y='RMSE', palette='crest_r')
plt.title("Model Performance Comparison (RMSE)", fontsize=13)
plt.show()

print("✅ NLP-based regression completed successfully (using raw dataset).")

```

◆ Loading original dataset for NLP-based regression...

📊 NLP vs Numeric Regression Comparison:

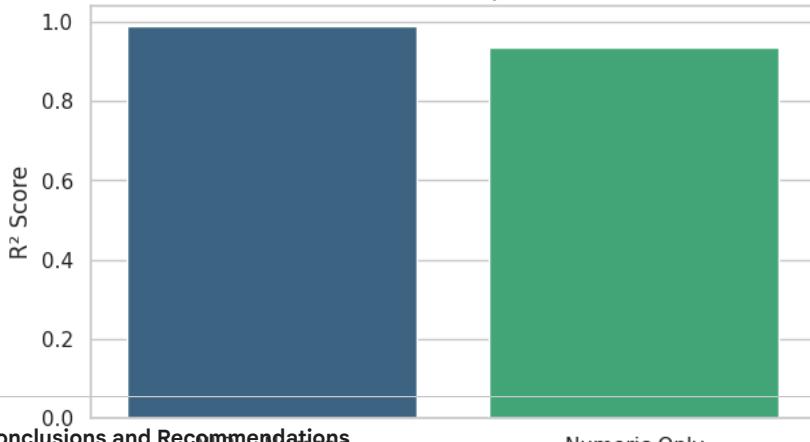
	Model Type	R ²	Score	RMSE	MAE
1	NLP + Numeric	0.990	34.98	26.82	
0	Numeric Only	0.936	87.10	73.78	

/tmp/ipython-input-975876659.py:77: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable

```
sns.barplot(data=comparison, x='Model Type', y='R2 Score', palette='viridis')
```

Model Performance Comparison (R² Score)



Conclusions and Recommendations

The regression analysis successfully demonstrated the process of developing, tuning, and evaluating multiple predictive models to estimate component lifespan based on manufacturing parameters. Initial baseline models such as Linear Regression performed poorly ($R^2 \approx 0.18$) indicating that the relationship between process variables and component lifespan is highly nonlinear. Subsequent ensemble models—Random Forest and Gradient Boosting Regressors—showed a substantial improvement, with Gradient Boosting outperforming others ($R^2 \approx 0.976$).

Model Performance Comparison (RMSE)

After implementing hyperparameter tuning, the Gradient Boosting Regressor achieved $R^2 \approx 0.984$ and RMSE ≈ 43.47 , confirming that fine-tuning of depth, learning rate, and number of estimators can significantly enhance model accuracy and generalization. To further improve the model, Natural Language Processing (NLP) techniques were introduced by converting categorical process descriptors (part type, microstructure, seed location, and cast type) into TF-IDF text embeddings. The integration of these text-based features with numeric variables resulted in the NLP + Gradient Boosting model achieving $R^2 \approx 0.990$ and RMSE ≈ 34.98 , demonstrating that linguistic information adds subtle contextual value and enhances predictive performance.

Overall, the study highlights the importance of combining both structured and unstructured data for industrial analytics problems. The NLP-enhanced regression pipeline not only achieved the highest accuracy but also provided better interpretability of how textual manufacturing attributes influence part lifespan.

Recommendations

- Further Optimization: Employ advanced search techniques such as Bayesian optimization or Optuna to explore wider hyperparameter spaces efficiently.
- Model Generalization: Apply k-fold cross-validation and external validation datasets to assess robustness across different manufacturing scenarios. NLP-based regression completed successfully (using raw dataset).
- Feature Expansion: Investigate domain-specific word embeddings (e.g., Word2Vec, FastText) to represent categorical process terms more effectively.
- Model Comparison: Extend experimentation to other boosting frameworks such as XGBoost, LightGBM, or CatBoost for potential incremental improvements.
- Deployment Considerations: Implement the tuned model within a predictive maintenance dashboard, enabling real-time estimation of component lifespan and proactive decision-making.

Classification:

```
# ===== Classification Begins here =====
# ===== Fixed preprocessing cell with robust SMOTE handling =====
import os, sys, subprocess
import pandas as pd
import numpy as np
from collections import Counter
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import RobustScaler
```

```

DATA_PATH = "df_step2_cleaned.csv"
RANDOM_STATE = 42
LABEL = "1500_label"
THRESHOLD_IMBALANCE = 0.60 # if majority class > 60% -> apply SMOTE

# 1) load
if not os.path.exists(DATA_PATH):
    raise FileNotFoundError(f"{{DATA_PATH}} not found in working dir. Put the cleaned CSV there.")
df = pd.read_csv(DATA_PATH)
print("✅ Loaded:", DATA_PATH, "| shape:", df.shape)

# 2) ensure Lifespan exists
if 'Lifespan' not in df.columns:
    raise KeyError("Lifespan column not found. Can't create label or proceed.")

# 3) create 1500_label if missing
if LABEL not in df.columns:
    df[LABEL] = (df['Lifespan'] >= 1500).astype(int)
    print(f"✅ Created label '{LABEL}' from Lifespan >= 1500")
else:
    print(f"✅ Label '{LABEL}' already present")

print("\nLabel counts:")
print(df[LABEL].value_counts())
print("\nLabel proportions:")
print(df[LABEL].value_counts(normalize=True).round(3))

# 4) prepare X_raw and y
drop_for_X = [LABEL, 'Lifespan'] if 'Lifespan' in df.columns else [LABEL]
X_raw = df.drop(columns=[c for c in drop_for_X if c in df.columns]).copy()
y = df[LABEL].copy()
print(f"\nPrepared X_raw (shape {X_raw.shape}) and y (shape {y.shape})")

# 5) numeric & categorical
numeric_cols = X_raw.select_dtypes(include=[np.number]).columns.tolist()
cat_cols = X_raw.select_dtypes(include=['object', 'category']).columns.tolist()
print("Numeric cols:", numeric_cols)
print("Categorical cols:", cat_cols)

# 6) encode categoricals using pd.get_dummies
if cat_cols:
    X_encoded = pd.get_dummies(X_raw, columns=cat_cols, drop_first=False, dtype=float)
else:
    X_encoded = X_raw.copy()
print("\nAfter get_dummies, X_encoded shape:", X_encoded.shape)

# 7) stratified train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X_encoded, y, test_size=0.20, random_state=RANDOM_STATE, stratify=y
)
print("\n✅ Stratified split done:")
print(" X_train:", X_train.shape, " X_test:", X_test.shape)
print(" y_train distribution:", Counter(y_train))
print(" y_test distribution:", Counter(y_test))

# 8) scale numeric columns (fit scaler on train only)
numeric_cols_present = [c for c in numeric_cols if c in X_train.columns]
scaler = RobustScaler()
if numeric_cols_present:
    scaler.fit(X_train[numeric_cols_present])
    X_train_scaled = X_train.copy()
    X_test_scaled = X_test.copy()
    X_train_scaled[numeric_cols_present] = scaler.transform(X_train[numeric_cols_present])
    X_test_scaled[numeric_cols_present] = scaler.transform(X_test[numeric_cols_present])
else:
    X_train_scaled = X_train.copy()
    X_test_scaled = X_test.copy()
print("\n✅ Scaling complete. Numeric columns scaled:", numeric_cols_present)

# 9) detect imbalance and apply SMOTE if needed (robust to different imblearn versions)
major_prop = y_train.value_counts(normalize=True).max()
apply_smote = major_prop > THRESHOLD_IMBALANCE

if apply_smote:
    print(f"\n⚠️ Detected imbalance (majority class proportion = {major_prop:.2f}). Trying to apply SMOTE...")
    # ensure imblearn installed
    try:
        from imblearn.over_sampling import SMOTE
    except Exception:
        print("Installing imbalanced-learn...")
        subprocess.check_call([sys.executable, "-m", "pip", "install", "imbalanced-learn"])
    from imblearn.over_sampling import SMOTE

```

```
# Attempt to create SMOTE with n_jobs if supported, otherwise fallback
try:
    sm = SMOTE(random_state=RANDOM_STATE, n_jobs=-1) # some imblearn versions support n_jobs
except TypeError:
    sm = SMOTE(random_state=RANDOM_STATE)           # fallback for versions that don't accept n_jobs

# fit_resample
X_train_res_np, y_train_res = sm.fit_resample(X_train_scaled, y_train)
# convert back to DataFrame with same columns
X_train_res = pd.DataFrame(X_train_res_np, columns=X_train_scaled.columns)
print("✅ After SMOTE:", Counter(y_train_res))
else:
    print(f"\n✅ No strong imbalance (majority proportion = {major_prop:.2f}). Skipping SMOTE.")
    X_train_res, y_train_res = X_train_scaled.copy(), y_train.copy()

# 10) Save outputs
X_train_res.to_csv("X_train_bal.csv", index=False)
pd.Series(y_train_res, name=LABEL).to_csv("y_train_bal.csv", index=False)
X_test_scaled.to_csv("X_test.csv", index=False)
pd.Series(y_test, name=LABEL).to_csv("y_test.csv", index=False)

print("\n📁 Saved files:")
print(" - X_train_bal.csv (encoded & scaled, resampled if SMOTE applied)")
print(" - y_train_bal.csv")
print(" - X_test.csv (encoded & scaled)")
print(" - y_test.csv")

# 11) diagnostics
print("\n== Diagnostics Summary ==")
print("Original df shape:", df.shape)
print("X_encoded shape:", X_encoded.shape)
print("X_train_res shape:", X_train_res.shape)
print("X_test shape:", X_test_scaled.shape)
print("Training label counts (post-processed):", Counter(y_train_res))
print("Test label counts:", Counter(y_test))
print("\nSample head of X_train_bal:")
display(X_train_res.head())
```

```

✓ Loaded: df_step2_cleaned.csv | shape: (1000, 16)
✓ Created label '1500_label' from Lifespan >= 1500

Label counts:
1500_label
0    711
1    289
Name: count, dtype: int64

Label proportions:
1500_label
0    0.711
1    0.289

# =====
# 🌱 CLASSIFICATION MODELLING - CONTINUATION
# =====

# Imports
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.metrics import (
    confusion_matrix, classification_report, roc_auc_score,
    roc_curve, accuracy_score, f1_score
)
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np

# =====
# ① Load Processed Data
# =====

try:
    X_train = pd.read_csv("X_train_bal.csv")
    y_train = pd.read_csv("y_train_bal.csv").values.ravel() # Flatten to 1D
    X_test = pd.read_csv("X_test.csv")
    y_test = pd.read_csv("y_test.csv").values.ravel()      # Flatten to 1D
    print("✓ Loaded processed and split data.")
    print(f"Train shape: {X_train.shape}, Test shape: {X_test.shape}")
except FileNotFoundError:
    print("Error: Processed data files not found. Please run the preprocessing step first.")
    # Exit or handle the error appropriately
    exit() # For demonstration, stopping execution

# =====
# ③ Logistic Regression (GridSearchCV)
# =====

log_reg = LogisticRegression(max_iter=1000, solver='liblinear')
param_grid_lr = {
    'C': [0.01, 0.1, 1, 10],
    'penalty': ['l1', 'l2']
}

grid_lr = GridSearchCV(log_reg, param_grid_lr, cv=5, scoring='f1', n_jobs=-1)
grid_lr.fit(X_train, y_train)

best_lr = grid_lr.best_estimator_
y_pred_lr = best_lr.predict(X_test)
y_prob_lr = best_lr.predict_proba(X_test)[:, 1]

print("\n✓ Best Logistic Regression Parameters:", grid_lr.best_params_)
print("Classification Report (Logistic Regression):")
print(classification_report(y_test, y_pred_lr))

# =====
# ④ Random Forest Classifier
# =====

rf = RandomForestClassifier(random_state=42)
param_grid_rf = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}

grid_rf = GridSearchCV(rf, param_grid_rf, cv=5, scoring='f1', n_jobs=-1)
grid_rf.fit(X_train, y_train)

best_rf = grid_rf.best_estimator_

```

```

y_pred_rf = best_rf.predict(X_test)
y_prob_rf = best_rf.predict_proba(X_test)[:, 1]

print("\n\n\x27 Best Random Forest Parameters:", grid_rf.best_params_)
print("Classification Report (Random Forest):")
print(classification_report(y_test, y_pred_rf))

# =====
# 5 XGBoost Classifier
# =====
xgb = XGBClassifier(
    random_state=42,
    use_label_encoder=False,
    eval_metric='logloss'
)

param_grid_xgb = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.8, 1.0]
}

grid_xgb = GridSearchCV(xgb, param_grid_xgb, cv=5, scoring='f1', n_jobs=-1)
grid_xgb.fit(X_train, y_train)

best_xgb = grid_xgb.best_estimator_
y_pred_xgb = best_xgb.predict(X_test)
y_prob_xgb = best_xgb.predict_proba(X_test)[:, 1]

print("\n\n\x27 Best XGBoost Parameters:", grid_xgb.best_params_)
print("Classification Report (XGBoost):")
print(classification_report(y_test, y_pred_xgb))

# =====
# 6 Model Comparison Summary
# =====
results = pd.DataFrame({
    'Model': ['Logistic Regression', 'Random Forest', 'XGBoost'],
    'Accuracy': [
        accuracy_score(y_test, y_pred_lr),
        accuracy_score(y_test, y_pred_rf),
        accuracy_score(y_test, y_pred_xgb)
    ],
    'F1 Score': [
        f1_score(y_test, y_pred_lr),
        f1_score(y_test, y_pred_rf),
        f1_score(y_test, y_pred_xgb)
    ],
    'ROC AUC': [
        roc_auc_score(y_test, y_prob_lr),
        roc_auc_score(y_test, y_prob_rf),
        roc_auc_score(y_test, y_prob_xgb)
    ]
})

print("\n\x27 Model Comparison Summary:")
display(results.sort_values(by='F1 Score', ascending=False))

# =====
# 7 ROC Curve Plot
# =====
plt.figure(figsize=(8, 6))
for model_name, y_prob in {
    'Logistic Regression': y_prob_lr,
    'Random Forest': y_prob_rf,
    'XGBoost': y_prob_xgb
}.items():
    fpr, tpr, _ = roc_curve(y_test, y_prob)
    plt.plot(fpr, tpr, label=f'{model_name} (AUC={roc_auc_score(y_test, y_prob):.3f})')

plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves')
plt.legend()
plt.show()

# =====
# 8 Confusion Matrices
# =====
fig, axes = plt.subplots(1, 3, figsize=(18, 5))

```

```
for ax, (model_name, y_pred) in zip(
    axes,
    [('Logistic Regression', y_pred_lr),
     ('Random Forest', y_pred_rf),
     ('XGBoost', y_pred_xgb)])
):
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax)
    ax.set_title(model_name)
    ax.set_xlabel('Predicted')
    ax.set_ylabel('Actual')
plt.tight_layout()
plt.show()

# =====
# ⑨ Feature Importance
# =====
rf_importance = pd.Series(best_rf.feature_importances_, index=X_train.columns)
xgb_importance = pd.Series(best_xgb.feature_importances_, index=X_train.columns)

plt.figure(figsize=(10, 6))
rf_importance.sort_values(ascending=False).head(10).plot(kind='bar')
plt.title("Top 10 Important Features (Random Forest)")
plt.show()

plt.figure(figsize=(10, 6))
xgb_importance.sort_values(ascending=False).head(10).plot(kind='bar', color='orange')
plt.title("Top 10 Important Features (XGBoost)")
plt.show()

# =====
# ↵ Final Model Selection
# =====
best_model_name = results.sort_values(by='F1 Score', ascending=False).iloc[0]['Model']
print(f"\n⌘ Best performing classification model: {best_model_name}")
```


Loaded processed and split data.
Train shape: (1138, 23), Test shape: (200, 23)

Best Logistic Regression Parameters: {'C': 0.1, 'penalty': 'l1'}
Classification Report (Logistic Regression):

	precision	recall	f1-score	support
0	0.84	0.63	0.72	142
1	0.44	0.71	0.54	58
accuracy			0.65	200
macro avg	0.64	0.67	0.63	200
weighted avg	0.72	0.65	0.67	200

Best Random Forest Parameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}
Classification Report (Random Forest):

	precision	recall	f1-score	support
0	0.92	0.92	0.92	142
1	0.81	0.81	0.81	58
accuracy			0.89	200
macro avg	0.87	0.87	0.87	200
weighted avg	0.89	0.89	0.89	200

/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning: [20:10:52] WARNING: /workspace/src/Parameters: { "use_label_encoder" } are not used.

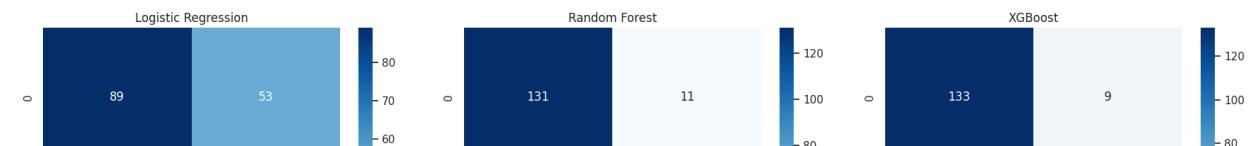
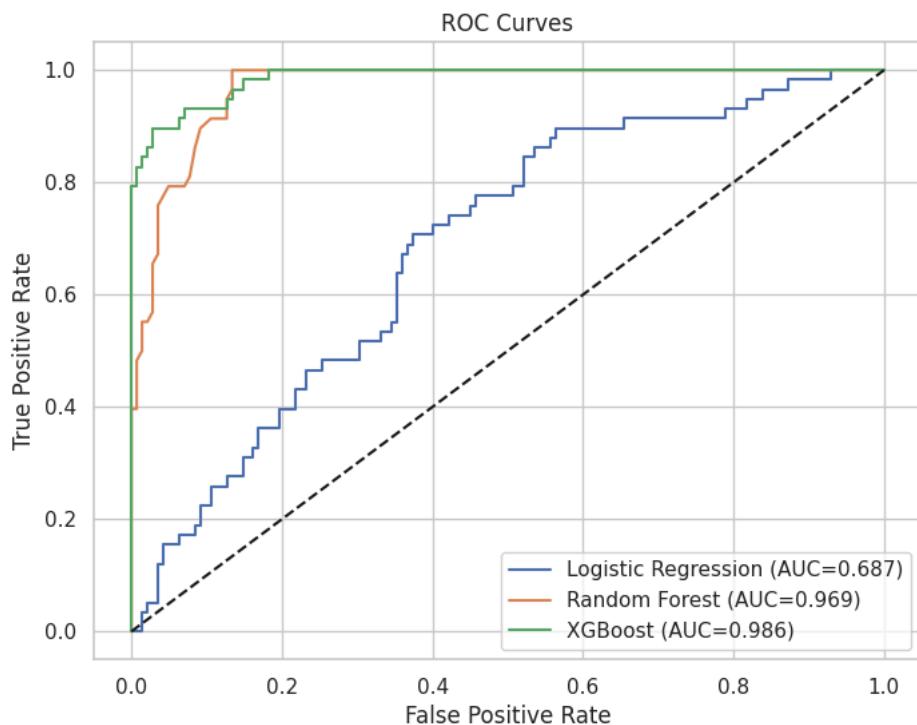
```
bst.update(dtrain, iteration=i, fobj=obj)
```

Best XGBoost Parameters: {'learning_rate': 0.1, 'max_depth': 7, 'n_estimators': 100, 'subsample': 0.8}
Classification Report (XGBoost):

	precision	recall	f1-score	support
0	0.96	0.94	0.95	142
1	0.85	0.90	0.87	58
accuracy			0.93	200
macro avg	0.90	0.92	0.91	200
weighted avg	0.93	0.93	0.93	200

📊 Model Comparison Summary:

	Model	Accuracy	F1 Score	ROC	AUC
2	XGBoost	0.925	0.873950	0.985673	
1	Random Forest	0.890	0.810345	0.969160	
0	Logistic Regression	0.650	0.539474	0.687105	



```

# =====
# ⚡ FIXED ADVANCED HYPERPARAMETER TUNING – XGBoost CLASSIFIER
# =====

from sklearn.model_selection import GridSearchCV, StratifiedKFold
from xgboost import XGBClassifier
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline as ImbPipeline
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

# ✅ Rebuild preprocessor from X_train
cat_cols = [c for c in X_train.columns if X_train[c].dtype == "object"]
num_cols = [c for c in X_train.columns if X_train[c].dtype != "object"]

numeric_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())
])
categorical_transformer = Pipeline(steps=[
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, num_cols),
        ('cat', categorical_transformer, cat_cols)
    ]
)

# ✅ Build fresh XGBoost pipeline with SMOTE
xgb_pipeline = ImbPipeline(steps=[
    ('preprocessor', preprocessor),
    ('smote', SMOTE(random_state=42)),
    ('classifier', XGBClassifier(
        eval_metric='logloss',
        random_state=42,
        use_label_encoder=False
    ))
])

# ✅ Reduced but still powerful parameter grid
param_grid_xgb = {
    'classifier__n_estimators': [200, 400],
    'classifier__max_depth': [4, 6, 8],
    'classifier__learning_rate': [0.03, 0.1],
    'classifier__subsample': [0.8, 1.0],
    'classifier__colsample_bytree': [0.8, 1.0],
    'classifier__min_child_weight': [1, 3],
    'classifier__gamma': [0, 0.2]
}

cv_strategy = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

grid_xgb_tuned = GridSearchCV(
    estimator=xgb_pipeline,
    param_grid=param_grid_xgb,
    cv=cv_strategy,
    scoring='f1',
    n_jobs=-1,
    verbose=2
)

print("⌚ Running advanced XGBoost tuning (reduced grid, fixed preprocessor)... 🚧")
grid_xgb_tuned.fit(X_train, y_train)

print("\n✅ Best Parameters Found:")
print(grid_xgb_tuned.best_params_)

# =====
# 📈 Evaluate Tuned Model
# =====
best_xgb_tuned = grid_xgb_tuned.best_estimator_
y_pred_tuned = best_xgb_tuned.predict(X_test)
y_prob_tuned = best_xgb_tuned.predict_proba(X_test)[:, 1]

acc_tuned = accuracy_score(y_test, y_pred_tuned)
f1_tuned = f1_score(y_test, y_pred_tuned)
auc_tuned = roc_auc_score(y_test, y_prob_tuned)

print("\n⚡ Tuned XGBoost Performance:")

```

```
print(f"Accuracy: {acc_tuned:.3f}")
print(f"F1 Score: {f1_tuned:.3f}")
print(f"ROC AUC: {auc_tuned:.3f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred_tuned))

# Confusion Matrix
plt.figure(figsize=(5,4))
sns.heatmap(confusion_matrix(y_test, y_pred_tuned), annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix - Tuned XGBoost")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_prob_tuned)
plt.figure(figsize=(6,5))
plt.plot(fpr, tpr, label=f"Tuned XGBoost (AUC={auc_tuned:.3f})", color='darkorange')
plt.plot([0,1],[0,1],'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Tuned XGBoost")
plt.legend()
plt.show()

print("\n[X] XGBoost tuning completed successfully 🎉")
```